

CPSC 8430 – Assignment 1

Part 1 Deep vs. Shallow

Task 1: Simulate a Function

- Two non-linear single-input, single-output functions were modeled for the first task: Y is equal to $\cos 1.5x$ and $\sin 2x$. Model 1, Model 2, and Model 3 were used to train the functions, and their parameters, 961, 945, and 975, were nearly identical to one another.
- The learning rate for each of the three models was 0.001. Figure 1 and Figure 2 depict the loss for Function 1 for Model 1, Model 2, and Model 3, respectively.

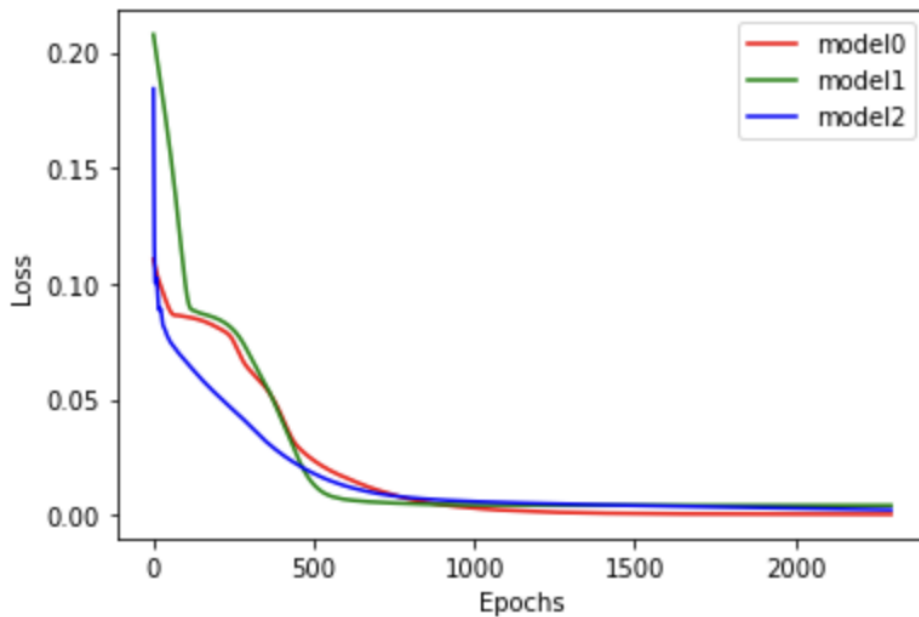


Figure 1

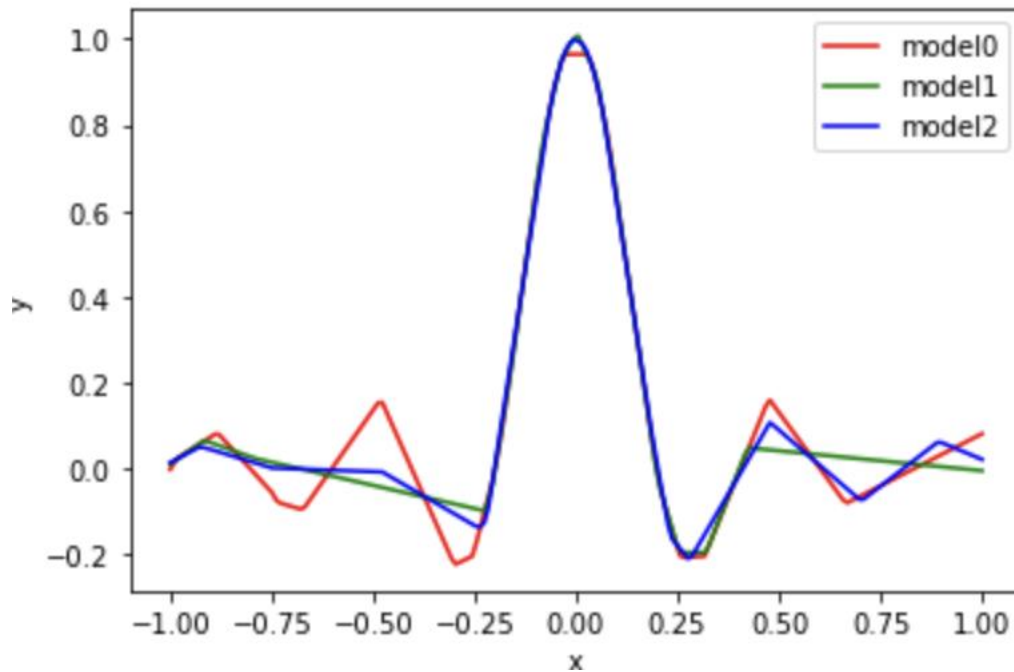


Figure 2

- For the first function, convergence was achieved by all three models around 2300 epochs ago, and for the second function, convergence was also achieved by all three models around 2300 epochs ago.
- The predicted and ground truth outcomes for the function $y = \sin 2x$ are depicted in Figure 3. The output of all the models is very similar, with a few exceptions. However, when input values range from 5.0 to 10.0, Models 2 and 3 seem to produce unusual output values.

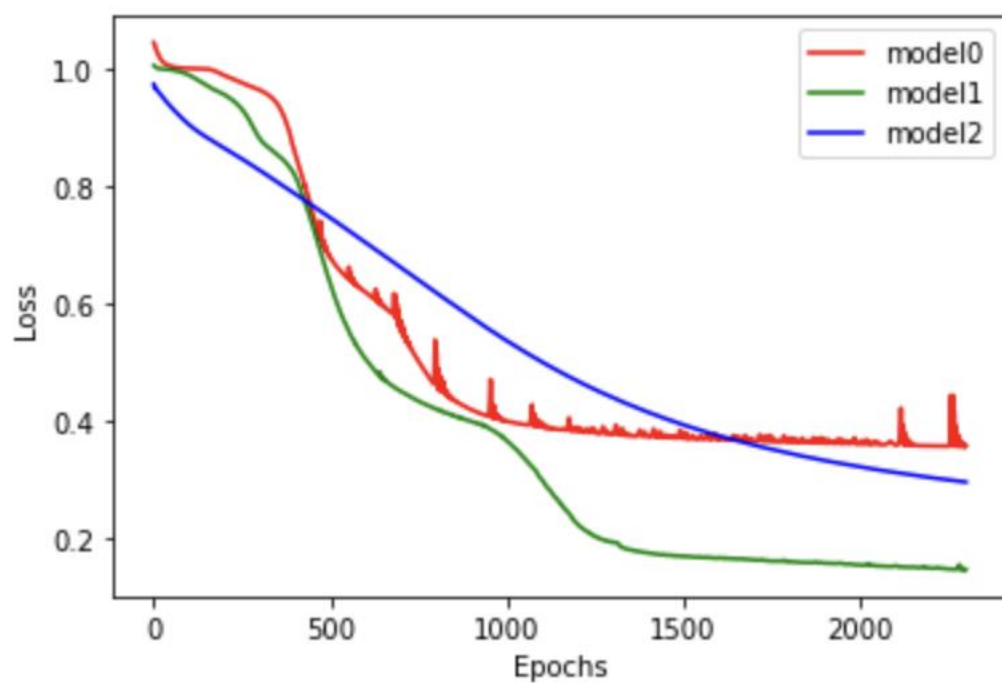


Figure 3

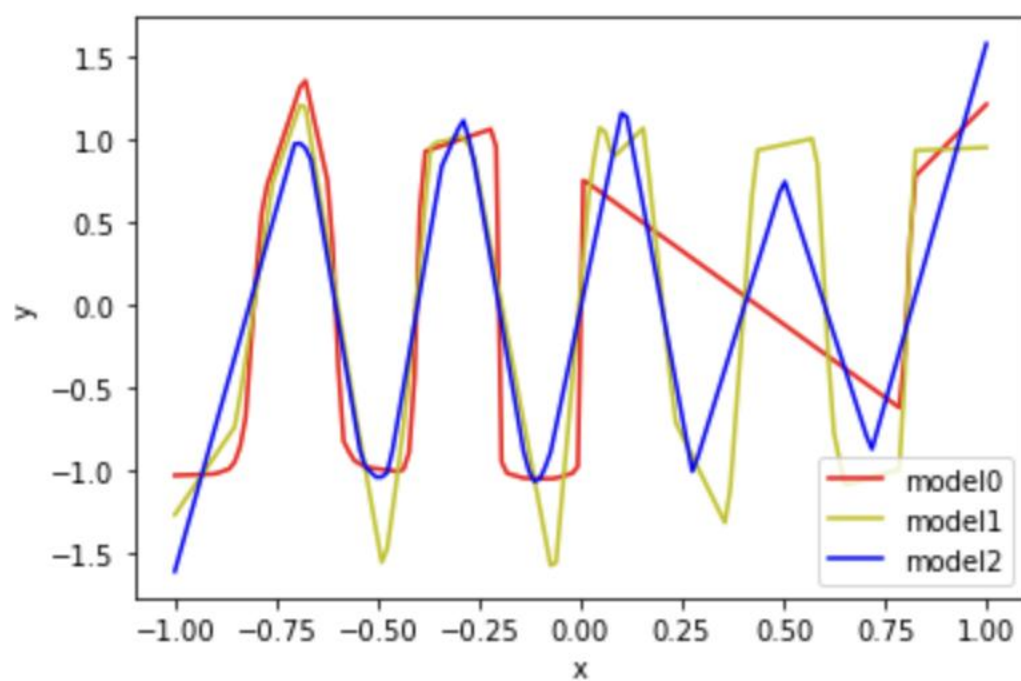


Figure 4

Task 2: Train on actual tasks

- The MNIST dataset was used to train three models for this task. The presentation of loss and accuracy charts served to illustrate the training process.
- Maximum pooling was utilized for pooling, and both models were fully connected. RELU was the activation function used. Each network had a different number of nodes in each layer. The learning rate for each model was 0.001. The loss that was sustained while model 1 and model 2 were being trained is depicted in Figure 5. The two models generally converged within 50 epochs, according to observations.

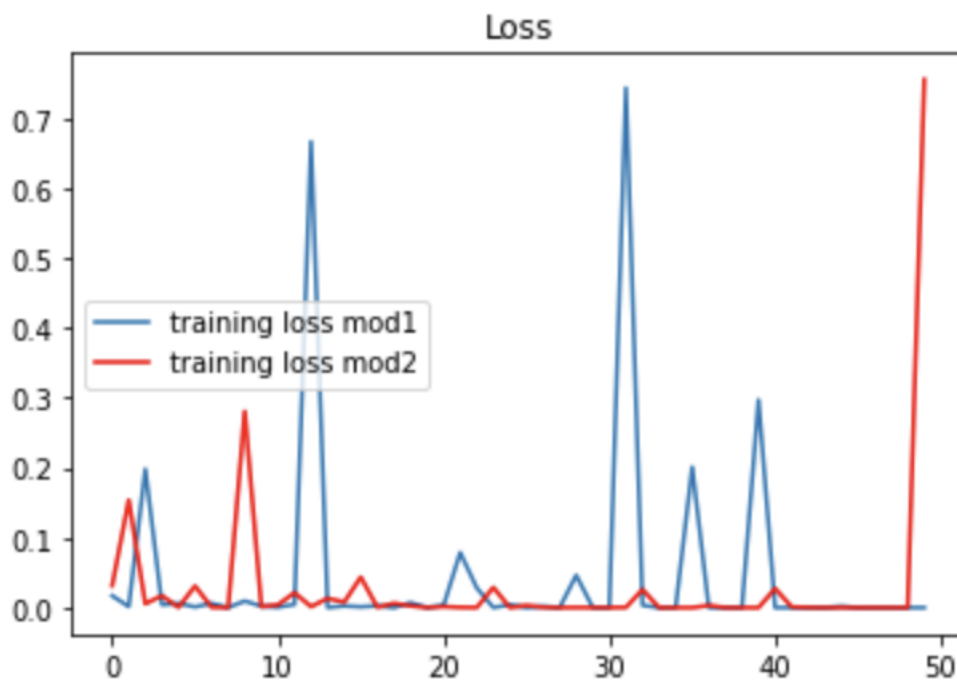


Figure 5

- The accuracy of the two models on training and test sets is shown in Figure 6 during model training. As expected, we can see that the two models regularly outperform each other on training data compared to test data.

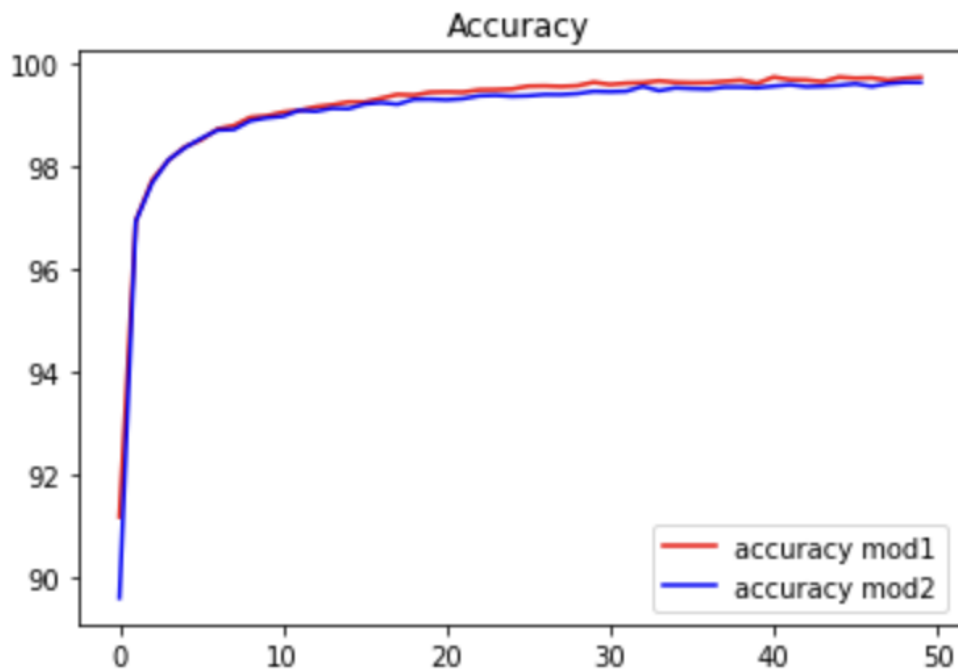


Figure 6

PART 2 – Optimization

Task 1 – Visualizing the Optimization process

- The modeled function was $y = \sin(2x)$. Figures 7 and 8 show how the network's weights were adjusted during the optimization process
- Backpropagation was used to improve the network's weights as it learned during the training phase. In the two mentioned figures, you can see how the weights have been gradually adjusted over time.

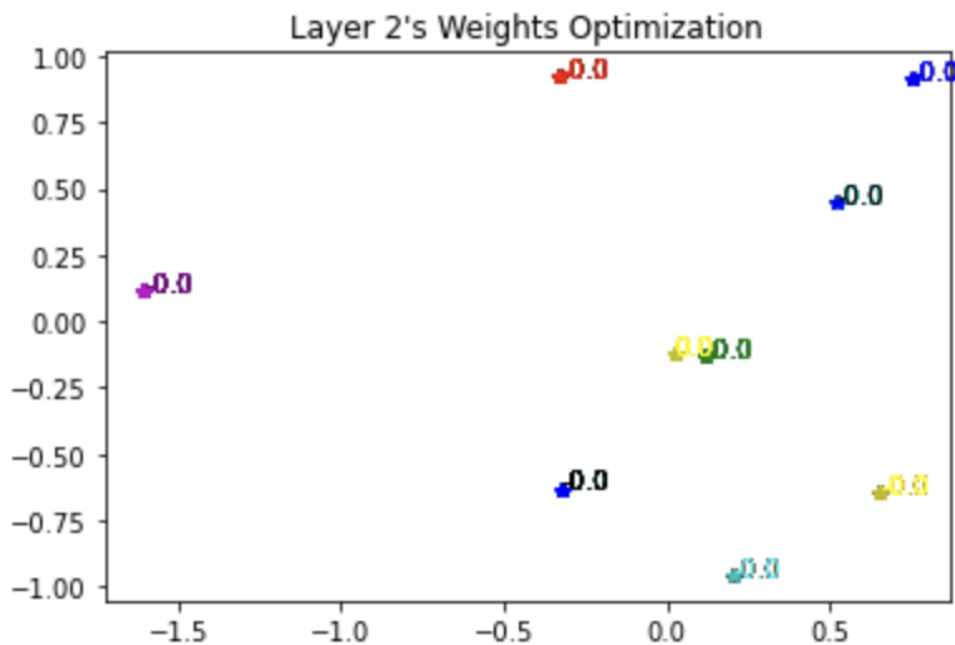


Figure 7

- The optimization of Layer 2's weights is depicted in Figure 7, while the optimization of the model's weights as a whole is depicted in Figure 8.

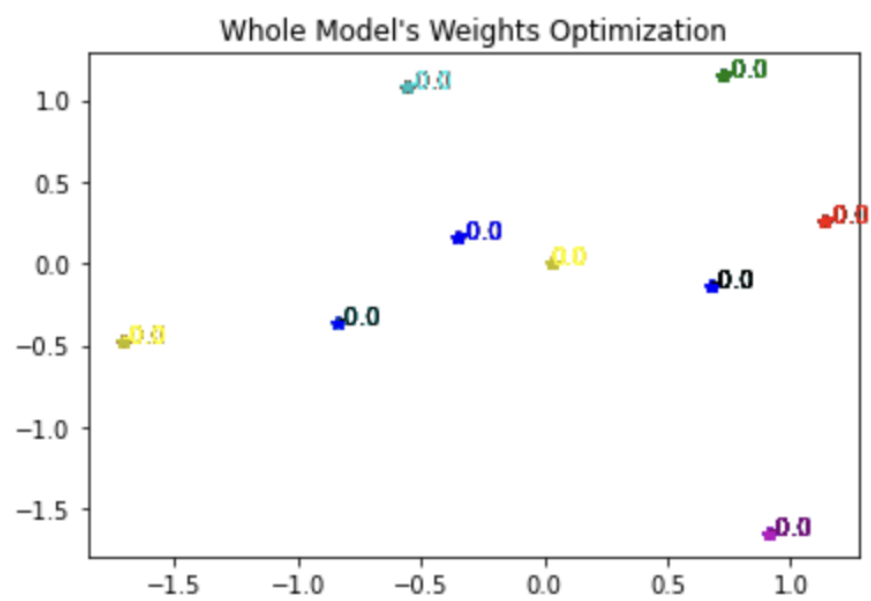


Figure 8

Task 2 – Observation of Gradient norm during training

- The model's loss is shown in Figure 9 for each iteration of training, while the gradient norm is shown in Figure 10. Each training epoch's loss for the model is depicted in Figure 9.
- The gradient norm for each epoch is shown in figure 10. Each spike in Figure 10 corresponds to a change in slope in Figure 9. The graph 10 differs from the ordinary.

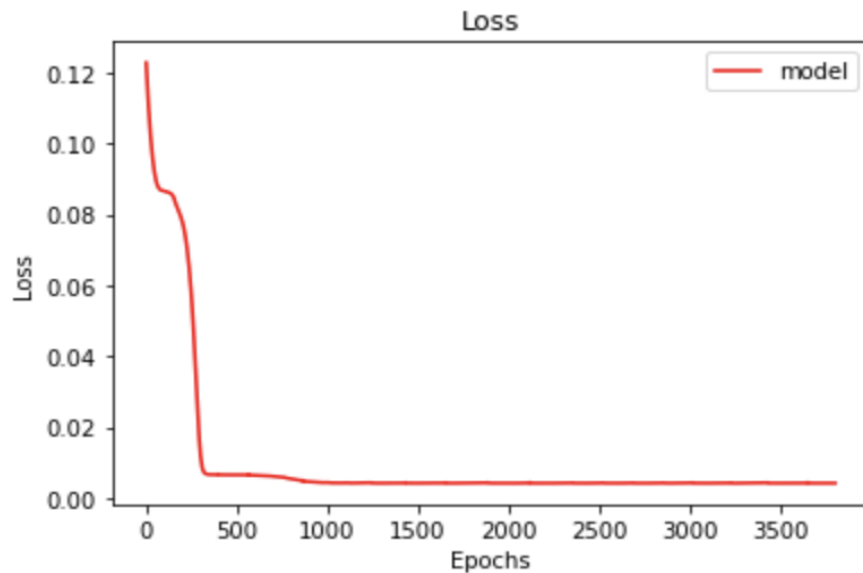


Figure 9

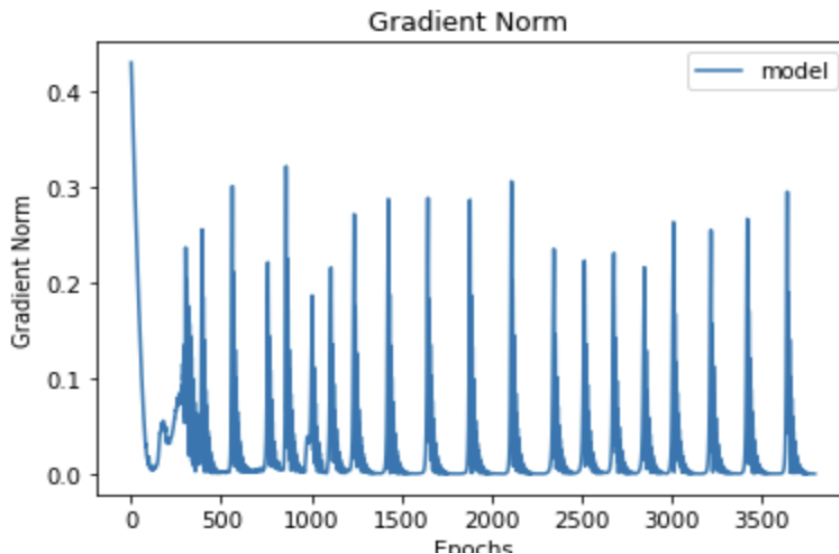


Figure 10

Part 3: Generalization

Task 1 – Can Network fit random labels?

- The MNIST dataset was used in this task to test the network's ability to fit random labels. For this purpose, a deep neural network model with three hidden layers was trained 3800 times on the MNIST dataset.
- The Adam optimizer was selected as the optimization algorithm and the learning rate was set to 0.001.

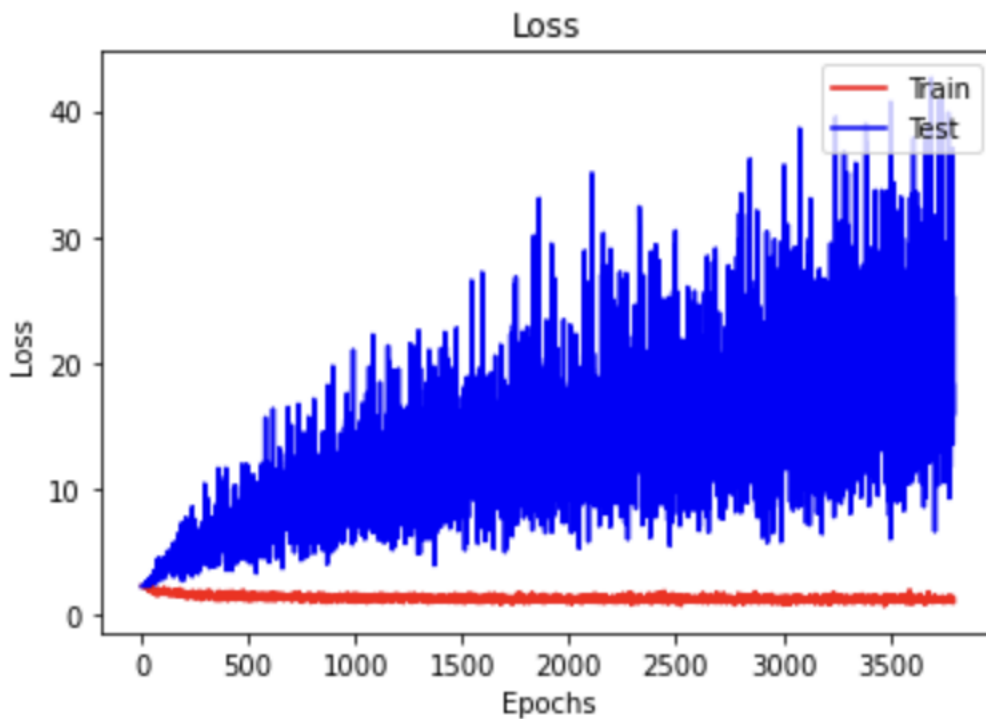


Figure 11

Task 2 – Number of Parameters vs. Generalization

- Ten distinct DNN models with various parameters ranging within specified values were trained using the MNIST dataset.
- The training data outperformed the test data in terms of accuracy and loss, as was to be expected. Even though the parameters were close to one million, there was always a difference between the test data and the training data.
- if we notice Figure 12 and 13, expanding the quantity of boundaries
- in the model abatements its misfortune and furthermore increment its precision. Be that as it may, after specific number of boundaries in a model, the model improvement from one cycle to another is unimportant adding extra boundaries further develops the model scarcely.

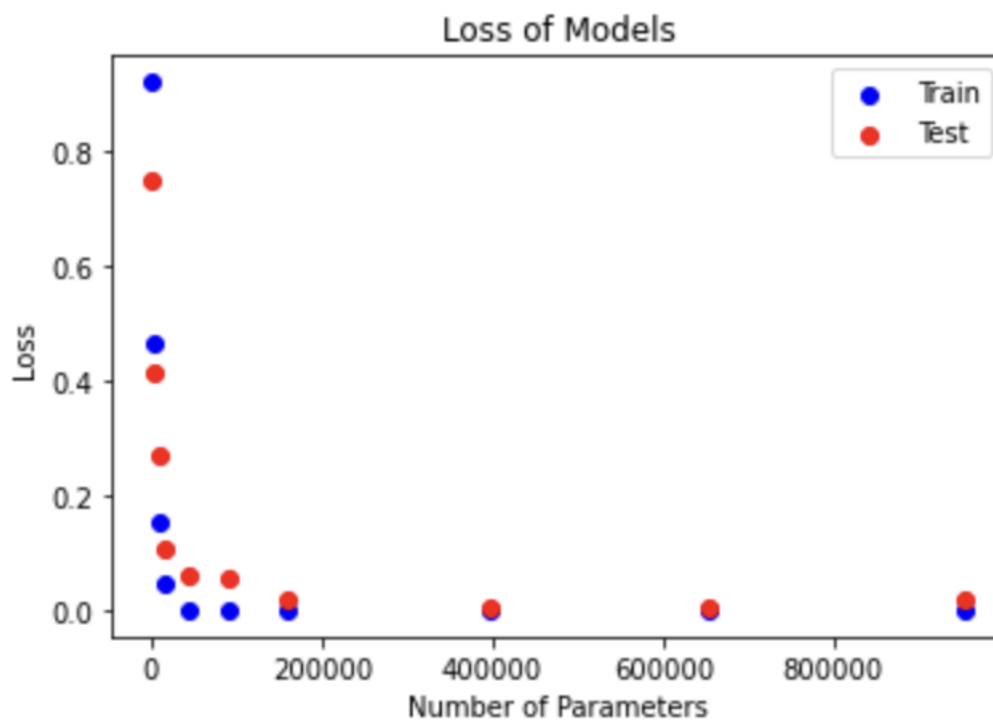


Figure 12

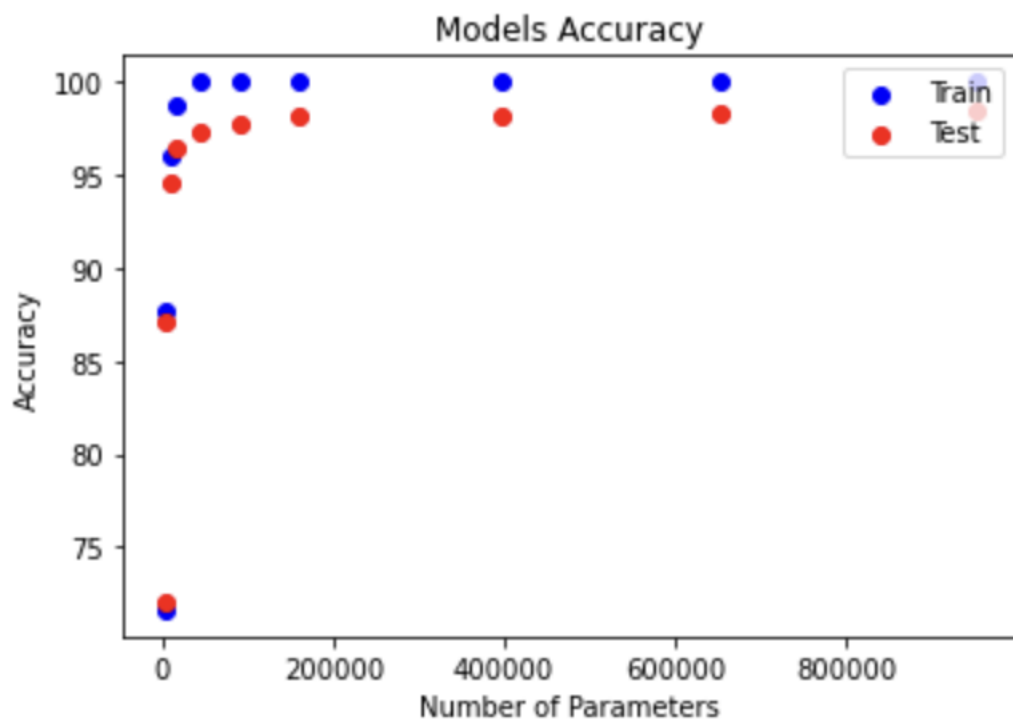


Figure 13

Task 3 – Flatness vs. Generalization

Part 1: -

- For the purpose of training and testing, the MNIST dataset was selected. Two separate batch sizes, 64 and 1024, were utilized to construct the two Deep Neural Networks. A learning rate of 0.001 was established for the entire model. The Adam Optimizer was utilized to optimize the neural network. Alpha represents the linear combination between θ_1 and θ_2 , with θ_1 representing the parameters of the first model and θ_2 representing the parameters of the second model.

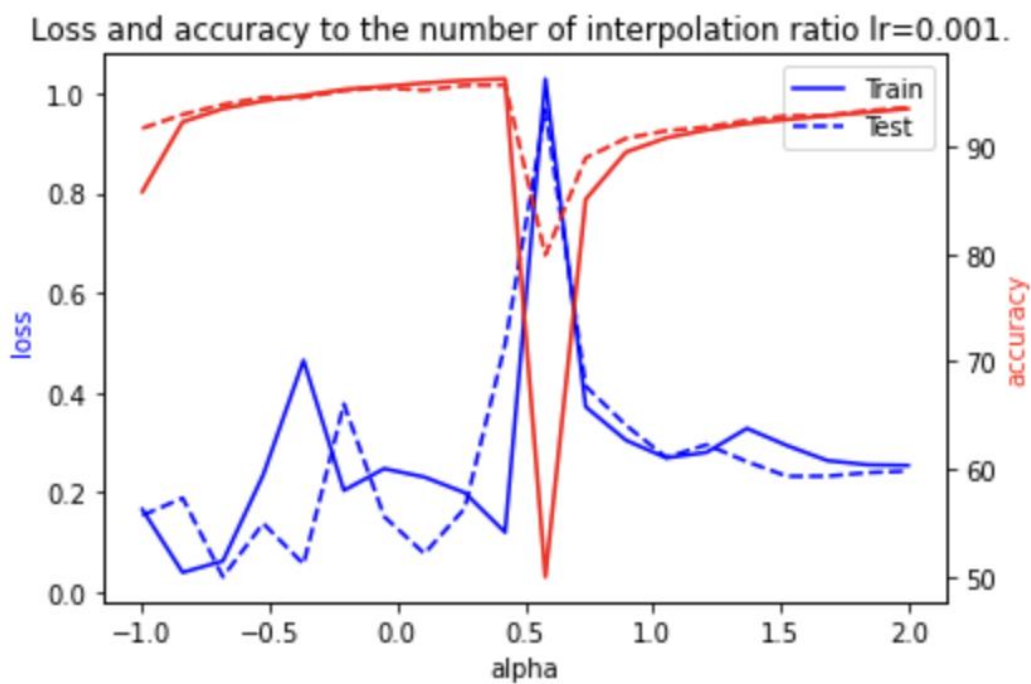


Figure 14

- With a learning rate of 0.001, the loss, accuracy, and linear interpolation alpha of the two models during training are shown in Figure 14. Figure 15, on the other hand, shows the same metrics but with a learning rate of 0.01.

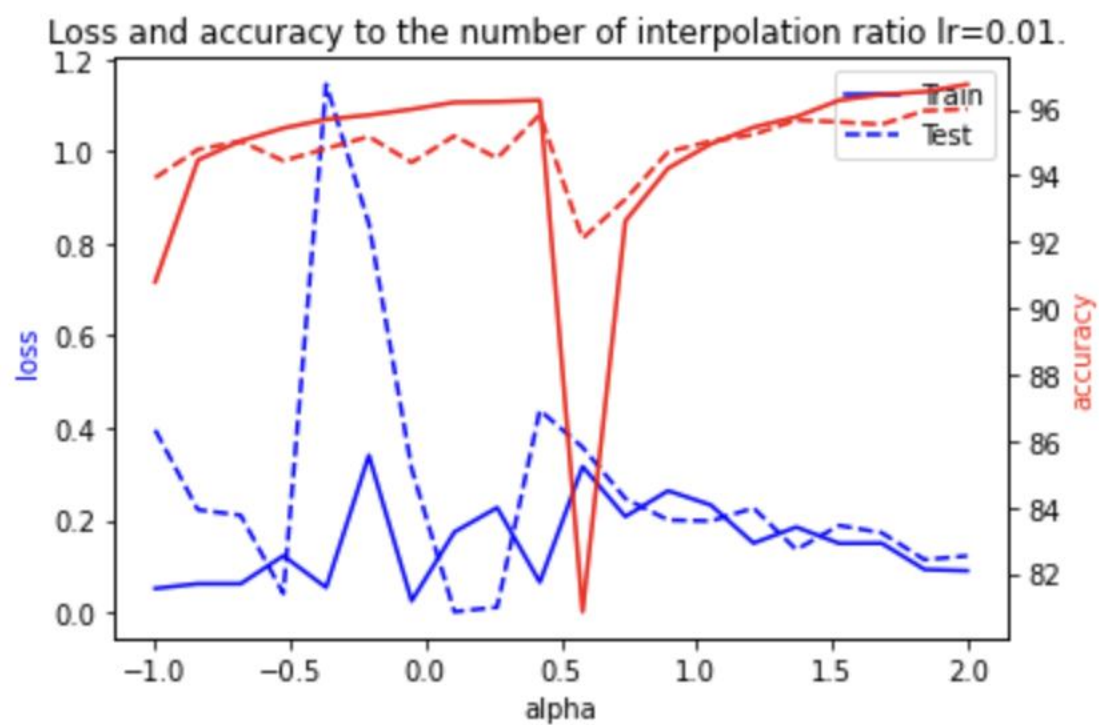


Figure 15

Part 2: Flatness Vs Generalization:

- The training and evaluation of the module were based on the MNIST data set. Each of the five similar Deep Neural Networks had two hidden layers and a total of 16630 parameters. Batches of 5 to 1000 were used to train these networks. In order to improve the optimization process, the Adam Optimizer was used, and each model had a learning rate of 0.001.

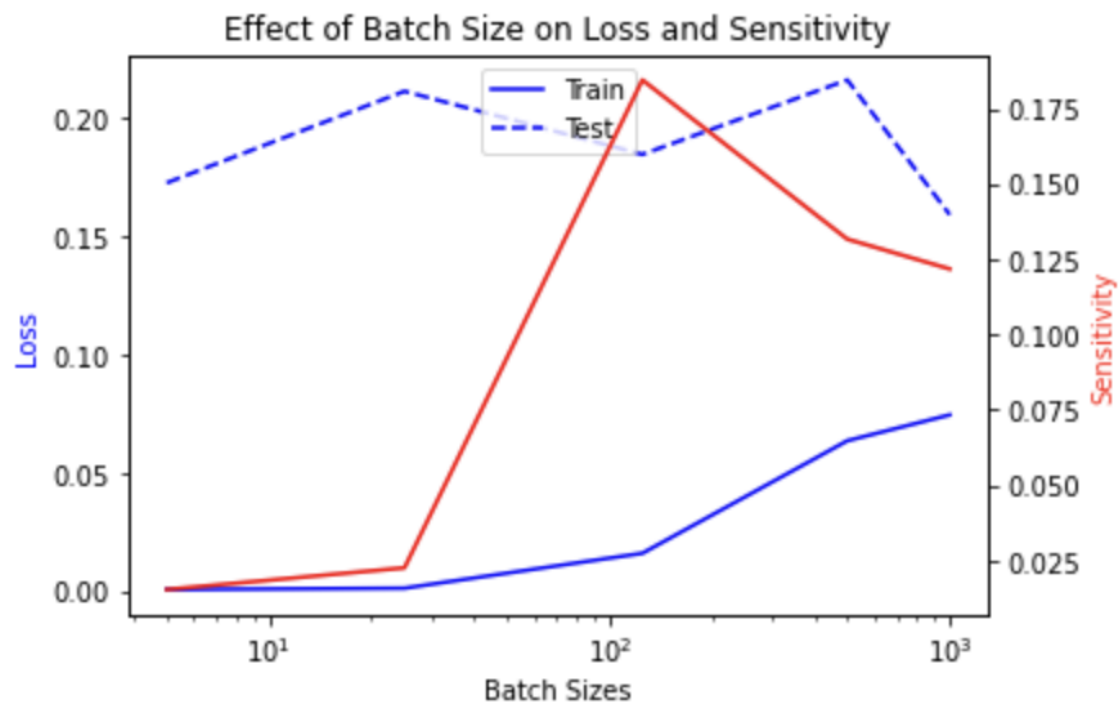


Figure 16

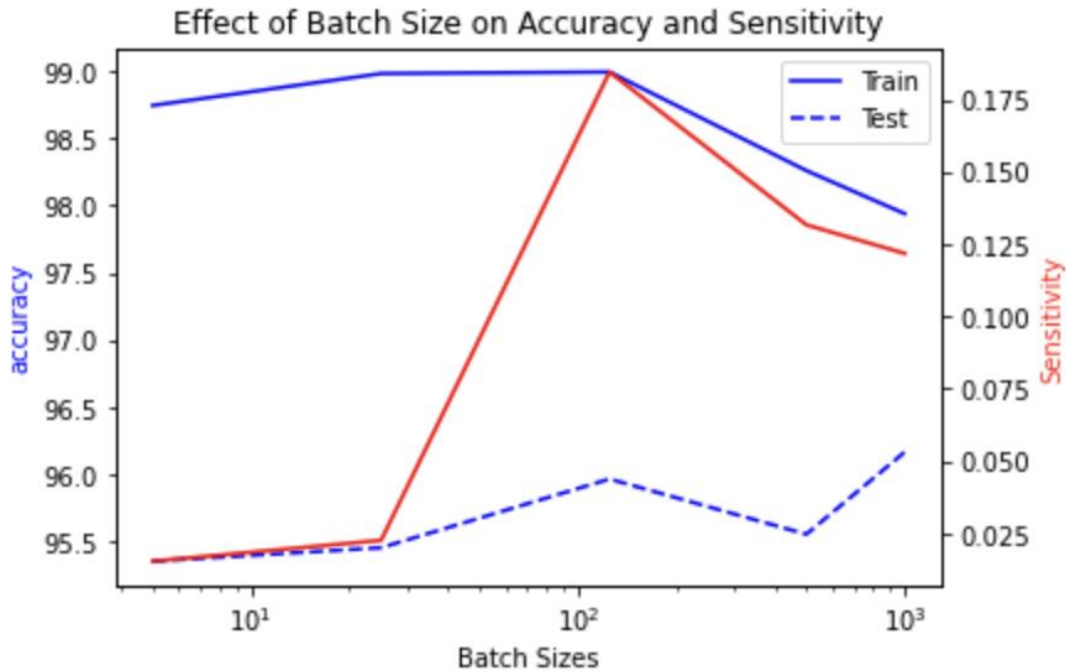


Figure 17

- The accuracy and loss of each of the five models were determined for the training dataset and the test dataset after the training process was completed. The Frobenius norm of the gradient method was then used to determine the models' sensitivity.
- Figures 16 and 17 show how batch size affects sensitivity and accuracy. Sensitivity decreases with increasing batch size. The network is most likely to produce the best results when the batch size is between 100 and 1000, as evidenced by this observation.

Github link- <https://github.com/Sujith-sai/DeeplearningHW1>